


Mira Project Screenshots: Workflow, In Action & Observability


The project consists of the n8n project workflows, the Slack workspace it posts into, the Drive folder it reads from, and one of the two live Google Sheets it reasons over and the business logic of ingesting files into pinecone vector DB and retrieving the details as required by the respective agents build as part of the project workflows to enable users with their queries.

Resource Files


Add resource files in the relevant Google Folder





Slack Configuration





Go to api.slack.com/apps → Create New App → From scratch. Name it after your project (e.g. 'Mira Bot') and select your workspace.




Under OAuth & Permissions, add the Bot Token Scopes: **chat:write**, **chat:write.public**, and **channels:read**, **users:read**, **users:read.email**. Click 'Install to Workspace' and authorize.



So once the permissions are added.. you need to install the App



Copy the generated Bot User OAuth Token (starts with xoxb-). In n8n, go to Credentials → New → Slack API, and paste it in. Invite the bot to your channel: in Slack, open #mira-updates and type /invite @YourBotName.




Workflow Canvas

Summary

Layer	Nodes	Purpose
Entry	N1 (chat trigger) / N1: Webhook, N2: User Input	Normalizes chat-widget and raw webhook input into one shape; tags source (chat vs webhook) so the final response routes back correctly
Orchestration	N6: Classifier Agent, N6: Simple Memory, N6a, N7, N8, N9	Classifies intent (plan / risk / status / unclear), validates the JSON, explodes multi-intent, routes
Specialist agents	N11 Planner, N16 Risk Assessor, N25 Status Reporter, M2 Milestone Tracker	Each is an @n8n/n8n-nodes-langchain.agent node grounded via a dedicated Pinecone retrieve-as-tool node, backed by gpt-4o-mini
Deterministic guardrails	N17d, N26b, N26d, plus every parser Code node	Non-LLM JavaScript that re-verifies, dedupes, re-buckets, or caps the LLM's own JSON output against the live Google Sheets data before it reaches the user — added because the LLM's own gatekeeping/bucketing proved unreliable in testing
Human-in-the-loop	N21b: New Risk Approval	Slack native sendAndWait double-approval (Approve/Deny) gate before any new risk is written to the register
Persistence	N17b/N22 (Risk Register), N26c (Task Board) — Google Sheets; Pinecone (mira-pineconeindex) — vector store	Live system-of-record reads/writes, separate from the RAG index
Output	Final Gate, Respond to Webhook, Post Message to Slack	Every response is posted to Slack channel mira-updates; webhook-sourced requests additionally get a synchronous HTTP response
Observability	LangSmith tracing (env vars) + evaluators	Full request/response tracing plus 4 offline evaluators (Hallucination, Assertions, Answer Relevance, Bias & Fairness)
Extended capability	Fine-tuning of the Risk Assessor's base model on corrected bug/fix examples	Optional enhancement per brief; see \$Fine-Tuning

Companion screenshots to the Workflow Architecture & Node Reference, PRD, and HLD docs — covering the n8n workflow canvas itself, Mira responding to a real request, and LangSmith observability.

Ingestion pipeline



Record count matches in pinecone index: 57 records

Pinecone / reachmesrividyap@gmail.com's Org / Default / Database

Docs Settings Get help R

PROJECT

- Overview
- Database
- Quickstart
- Indexes (3)
- Backups
- Assistant
- Inference

MANAGE

- API keys
- Settings

Database Usage

- RU/s 2 / 1M
- Storage 1.55MB / 20B
- WUs 2.6K / 2M
- Upgrade now

mira-pineconeindex

Record count 57 Region us-east-1 Model llama-text-embed-v2 Host mira-pineconeindex-uvy3gm.svc.apes-4627-574a...

BROWSER METRICS NAMESPACES (1) IMPORTS CONFIGURATION

Records

Namespace: NS_miraprojectplanningfiles

Operation: Search by text

Text: Ex: 'hello world'

+ Filter + Rerank

Top K 10 Search

Search: 10 results (top_k=10) 7 embed tokens 1 RUs 4.5 kB egress

1

SCORE 0.0632

_id: 5f29d35f-e1f1-4c30-a8fd-00038da7557e
blobType: "text/csv"
line: 6
loc.lines.from: 1
loc.lines.to: 7
source: "risk_register"
text: "risk_id: R06\ncategory: Cost Overruns\nrisk_challenge: Underestimation of AI development costs\nimpact: Budget con..."

2

SCORE 0.0613


_id: c71f1a3-cb72-4047-aac6-8ec57d3495c8
blobType: "text/csv"
line: 5
loc.lines.from: 1
loc.lines.to: 7
source: "risk_register"
text: "risk_id: R05\ncategory: Ethical & Legal Concerns\nrisk_challenge: Risks of biased models or data privacy violations\nl..."

3

SCORE 0.0581

_id: bd8d413c-0a9f-430f-ad21-6c1fd708d35f
blobType: "text/csv"
line: 20
loc.lines.from: 1

Retrieval pipeline




Mira in Action

Use case T1 with Chat Trigger

Generate a project plan for the AI Adoption Project at ABCDE Ltd. Use the provided project description.

ID	Test Input	Type	Expected Output Must Contain	Actual Output
T1	Generate a project plan for the AI Adoption Project at ABCDE Ltd. Use the provided project description.	Plan (detailed)	Must include phases, milestones, timeline. Must be grounded in the actual project description. Must NOT invent goals not in the description.	 The project plan is for the AI Adoption Project at ABCDE Ltd., which aims to enhance operational efficiency and improve customer experience through the integration of AI-powered tools. This plan is grounded in the project description and task documents retrieved, which outline specific phases, tasks, and key performance indicators such as reducing delivery delays by 15% and lowering operational costs by 10%.  Here's the project plan: #### Phase 1: Project Initiation (31-08-2026 to 08-09-2026) Activities: Define project scope and objectives, Form cross-functional task force Milestones: Project scope documented, Task force established Deliverables: Project scope document, Task force team list #### Phase 2: Data Infrastructure Assessment (31-08-2026 to 22-09-2026) Activities: Audit existing data infrastructure, Identify data gaps and quality issues Milestones: Data audit completed, Data gap analysis report Deliverables: Data infrastructure audit report, Gap analysis report #### Phase 3: Use Case Evaluation (28-09-2026 to 06-10-2026) Activities: Shortlist AI use cases, Vendor evaluation for ML platform, Legal review of data usage Milestones: Use cases shortlisted, Vendor evaluation completed Deliverables: Use case prioritization document, Vendor comparison report, Compliance review document #### Phase 4: Pilot Design (12-10-2026 to 08-11-2026) Activities: Establish data pipelines for AI models Milestones: Data pipelines established Deliverables: Data pipeline architecture document #### Phase 5: Pilot Implementation (09-11-2026 to 06-12-2026) Activities: Implement AI solutions in selected departments, Collect user feedback and performance metrics Milestones: User feedback and performance metrics collected Deliverables: Pilot implementation report, User feedback summary #### Phase 6: Staff Training (07-12-2026 to 25-01-2027) Activities: Develop and execute staff training program, Complete staff training rollout at remaining sites Milestones: Training program completed, Training extended to all sites Deliverables: Training materials, Training completion report This structured plan aligns with the project's goals and timelines, ensuring that all activities contribute to the overall success metrics of reducing delivery delays by 15% and lowering operational costs by 10%. Automated with this n8n workflow

Screenshots as below:




Use case T2 with Chat Trigger

Generate a project plan for: 'We want to build a chatbot.'

ID	Test Input	Type	Expected Output Must Contain	Actual Output
T2	Generate a project plan for: 'We want to build a chatbot.'	Plan (vague)	Must flag as insufficient detail. Must ask for scope, timeline, team size. Must NOT generate a full plan from one sentence.	I need a bit more detail to generate a plan. Here's what's missing: specific goals, timeline, and deliverables for the chatbot project. Could you share the project's scope and timeline? Automated with this n8n workflow

Screenshots as below:




Use case T3 with Chat Trigger

Generate a risk assessment for the AI Adoption Project at ABCDE Ltd.

ID	Test Input	Type	Expected Output Must Contain	Actual Output
T3	Generate a risk assessment for the AI Adoption Project at ABCDE Ltd.	Risk (detailed)	Must produce categorized risks relevant to logistics/AI. Must NOT include generic risks unrelated to project.	⚠️ Here are the relevant risks for the AI Adoption Project at ABCDE Ltd (all already tracked in the register). - Employees resist adopting AI-driven workflows [Change Resistance] Impact: Low adoption and reduced ROI; Mitigation: Strong change management and early engagement - Lack of internal AI/ML expertise [Talent & Skills Gap] Impact: Overreliance on external vendors and slow progress; Mitigation: Provide staff training and hire specialized talent - Difficulty integrating AI with legacy systems [Integration Issues] Impact: Delays in deployment and operational disruptions; Mitigation: Plan integration roadmap and use APIs/middleware - AI models lose accuracy over time [Model Performance Drift] Impact: Incorrect decisions and trust erosion; Mitigation: Regular retraining cycles and monitoring dashboards - Underestimation of AI development costs [Cost Overruns] Impact: Budget constraints and project delays; Mitigation: Set realistic budget and build in contingency funds - Risks of biased models or data privacy violations [Ethical & Legal Concerns] Impact: Legal repercussions and reputational damage; Mitigation: Involve legal early and implement fairness checks - AI models targeted by adversarial attacks [Security Vulnerabilities] Impact: Data breaches and operational risks; Mitigation: Adopt AI security best practices and penetration testing - Incomplete or unstructured data [Data Quality & Availability] Impact: Poor model performance and inaccurate outcomes; Mitigation: Perform early data audits and implement robust data governance Automated with this n8n workflow

Screenshots as below:



Use case T4 with Webhook

Generate a risk assessment for: 'New project starting soon.'

ID	Test Input	Type	Expected Output Must Contain	Actual Output
T4	Generate a risk assessment for: 'New project starting soon.'	Risk (vague)	Must flag as insufficient. Must ask for project details. Must NOT invent project-specific risks from nothing.	⚠️ I need more detail to assess risk properly. Here's what's missing: Specific project name or details about the new project.. Could you share more specifics for the details you are looking for? Automated with this n8n workflow

Screenshots as below:

HERA • PROJECT INTELLIGENCE ASSISTANT

Ask Mira

Submits to your n8n workflow's webhook. Works for plan, risk, and status requests, same as the chat window.

Webhook URL

http://localhost:5678/webhook/MLRobot-Newellined

Your request

Generate a risk assessment for: 'New project starting soon.'

Send request

Done.

I need more detail to assess risk properly.

Here's what's missing: Specific project name or details about the new project..

Could you share more specifics for the details you are looking for?

Automated with this [n8n workflow](#)

Channels

errors

mira-updates

n8nflowgoogledrive-projectupdates

project-questions-errors

I need more detail to assess risk properly.

Here's what's missing: Specific project name or details about the new project..

Could you share more specifics for the details you are looking for?

Automated with this [n8n workflow](#)

Use case T5 with Webhook

Generate a weekly status report using the sample task board for Sprint 3.

ID	Test Input	Type	Expected Output Must Contain	Actual Output
T5	Generate a weekly status report using the sample task board for Sprint 3.	Status (data)	Must show tasks by status. Must reference actual task names from CSV. Must NOT invent tasks.	Scoped to Sprint 3 — 4 tasks found, 1 milestone found. Status report for Sprint 3 (overall health: Green): To do: T009 - Vendor evaluation for ML platform, T025 - Legal review of data usage, T008 - Define success criteria for each use case, T007 - Shortlist AI use cases Upcoming milestones: M05 - Final prioritized use case list with success criteria Automated with this n8n workflow

Screenshots as below:

HERA • PROJECT INTELLIGENCE ASSISTANT

Ask Mira

Submits to your n8n workflow's webhook. Works for plan, risk, and status requests, same as the chat window.

Webhook URL

http://localhost:5678/webhook/MLRobot-Newellined

Your request

Generate a weekly status report using the sample task board for Sprint 3.

Send request

Done.

Scoped to Sprint 3 — 4 tasks found, 1 milestone found.

Status report for Sprint 3 (overall health: Green):

To do: T009 - Vendor evaluation for ML platform, T025 - Legal review of data usage, T008 - Define success criteria for each use case, T007 - Shortlist AI use cases

Upcoming milestones: M05 - Final prioritized use case list with success criteria

Drag and drop important stuff here

Channels

errors

mira-updates

n8nflowgoogledrive-projectupdates

project-questions-errors

Scoped to Sprint 3 — 4 tasks found, 1 milestone found.

Status report for Sprint 3 (overall health: Green):

To do: T009 - Vendor evaluation for ML platform, T025 - Legal review of data usage, T008 - Define success criteria for each use case, T007 - Shortlist AI use cases

Upcoming milestones: M05 - Final prioritized use case list with success criteria

Automated with this [n8n workflow](#)

Use case T6 with Webhook

Generate a weekly status report for: 'Things are going fine.'

ID	Test Input	Type	Expected Output Must Contain	Actual Output
T6	Generate a weekly status report for: 'Things are going fine.'	Status (filler, no genuine ask) — Assertions	Must ask the user to clarify what they'd like a status update on. Must NOT generate a full report (done/ in_progress/ to_do/ blocked/ at_risk/ upcoming_milestones must all be empty). Must NOT fabricate task statuses or progress percentages.	 Could you share more details on what you'd like a status update on? <i>Automated with this n8n workflow</i>

Screenshots as below:

MIRA - PROJECT INTELLIGENCE ASSISTANT

Ask Mira

Submits to your n8n workflow's webhook. Works for plan, risk, and status requests, same as the chat window.

Webhook URL

http://localhost:5678/webhook/MiraBot-!eueRef!ned

Your request

Generate a weekly status report for: 'Things are going fine.'

Send request

Done.

 Could you share more details on what you'd like a status update on?

mira-updates

n8nflowgoogledrive-projectupdates

project-questions-errors

 Could you share more details on what you'd like a status update on?

Automated with this [n8n workflow](#)

B I U     

Use case T7 with Webhook

What are the top 3 risks for the ABCDE Ltd project?

ID	Test Input	Type	Expected Output Must Contain	Actual Output
T7	What are the top 3 risks for the ABCDE Ltd project?	Risk analysis	Must reference actual risks from risk CSV. Must NOT invent risks not in the data.	 Here are the top 3 risks for the ABCDE Ltd project, ranked by severity. - Incomplete or unstructured data [Data Quality & Availability] Impact: Poor model performance and inaccurate outcomes; Mitigation: Perform early data audits and implement robust data governance - Difficulty integrating AI with legacy systems [Integration Issues] Impact: Delays in deployment and operational disruptions; Mitigation: Plan integration roadmap and use APIs/middleware - Lack of internal AI/ML expertise [Talent & Skills Gap] Impact: Overreliance on external vendors and slow progress; Mitigation: Provide staff training and hire specialized talent <i>Automated with this n8n workflow</i>

Screenshots as below:

MRSA - PROJECT INTELLIGENCE ASSISTANT

Ask Mira

Submits to your n8n workflow's webhook. Works for plan, risk, and status requests, same as the chat window.

Webhook URL

http://localhost:5678/webhook/!lnebot-llew@fined

Your request

What are the top 3 risks for the ABCDE Ltd project?

Send request

Done.

Here are the top 3 risks for the ABCDE Ltd project, ranked by severity.

- Incomplete or unstructured data [Data Quality & Availability] Impact: Poor model performance and inaccurate outcomes; Mitigation: Perform early data audits and implement robust data governance

- Difficulty integrating AI with legacy systems [Integration Issues] Impact: Delays in deployment and operational disruptions; Mitigation: Plan integration roadmap and use APIs/middleware

- Lack of internal AI/ML expertise [Talent & Skills Gap] Impact: Overreliance on external vendors and slow progress; Mitigation: Provide staff training and hire specialized talent

☆ Starred

Drag and drop important stuff here

Channels

errors

mira-updates

n8nflowgoogledrive-projectupdates

Here are the top 3 risks for the ABCDE Ltd project, ranked by severity.

- Incomplete or unstructured data [Data Quality & Availability] Impact: Poor model performance and inaccurate outcomes; Mitigation: Perform early data audits and implement robust data governance

- Difficulty integrating AI with legacy systems [Integration Issues] Impact: Delays in deployment and operational disruptions; Mitigation: Plan integration roadmap and use APIs/middleware

- Lack of internal AI/ML expertise [Talent & Skills Gap] Impact: Overreliance on external vendors and slow progress; Mitigation: Provide staff training and hire specialized talent

Automated with this n8n workflow

Use case T8 with Chat Trigger

Which tasks are blocked or at risk of missing their deadline?

ID	Test Input	Type	Expected Output Must Contain	Actual Output
T8	Which tasks are blocked or at risk of missing their deadline?	Tracking	Must identify T024 (Security review - BLOCKED). Must check due dates against current progress.	 Scoped to All Sprints — 1 task found.  Blocked & At risk — All Sprints: Blocked: T024 - Security review of AI infrastructure (waiting for security team availability) Automated with this n8n workflow

Screenshots as below:

The screenshot is divided into two panels. The left panel shows the n8n workflow editor with a complex flowchart of nodes and connections. A red box highlights the 'Execute workflow from N1: Workflow' button. Below the editor, a chat window shows a message: 'Which tasks are blocked or at risk of missing their deadline?' followed by a response: 'N33: Guard Empty Agent Output' and 'Final Gate'. The right panel shows a Slack message from 'Mira Bot' at 11:27 AM. The message content is: 'Scoped to All Sprints - 1 task found. Blocked & At risk - All Sprints: Blocked: T024 - Security review of AI infrastructure (waiting for security team availability) Automated with this n8n workflow'. Below the message is a rich text editor with various formatting options.

Use case T9 with Webhook

Generate a project plan for a 2-week project.

ID	Test Input	Type	Expected Output Must Contain	Actual Output
T9	Generate a project plan for a 2-	Plan (edge)	Must ask for scope, goals, deliverables. Must NOT generate detailed plan	 I need a bit more detail to generate a plan. Here's what's missing: specific project scope, goals, or deliverables.

	week project.		from timeline alone.	Could you share the project's scope and timeline? <i>Automated with this n8n workflow</i>
--	---------------	--	----------------------	--

Screenshots as below:

MIRA · PROJECT INTELLIGENCE ASSISTANT

Ask Mira

Submits to your n8n workflow's webhook. Works for plan, risk, and status requests, same as the chat window.

Webhook URL

http://localhost:5678/webhook/Mirabot-NewbieFined

Your request

Generate a project plan for a 2-week project.

Send request

Done.

I need a bit more detail to generate a plan.

Here's what's missing: specific project scope, goals, or deliverables.

Could you share the project's scope and timeline?

Project Intelligence Assistant

Direct messages

Agents & apps

Mira Bot

Project Updates

Connect apps

Mira Bot APP · 11:57 AM

I need a bit more detail to generate a plan.

Here's what's missing: specific project scope, goals, or deliverables.

Could you share the project's scope and timeline?

Automated with this [n8n workflow](#)

Use case T10 with Webhook

(This Example is shown in the Observability section on Tracing and Prompt fixes done further to improve the response)

Summarize current project status: tasks done, in progress, to do, blocked

ID	Test Input	Type	Expected Output Must Contain	Actual Output
T10	Summarize current project status: tasks done, in progress, to do, blocked?	Status summary (Observed Hallucination in Langsmith)	Must count correctly from task board CSV. Counts must match actual data.	 Scoped to All Sprints — 4 done, 23 to do, 1 blocked, 1 at risk, 12 milestones found.  Status report for All Sprints (overall health: Red): Done: T001 - Define project scope and objectives, T002 - Form cross-functional task force, T003 - Initial stakeholder alignment meeting, T004 - Audit existing data infrastructure To do: T006 - Identify data gaps and quality issues, T007 - Shortlist AI use cases, T008 - Define success criteria for each use case, T009 - Vendor evaluation for ML platform, T010 - Collect historical demand data, T011 - Data cleaning and preprocessing pipeline, T012 - Design demand forecasting model, T013 - Design route optimization model, T014 - Build model training pipeline, T015 - Deploy demand model to test environment, T016 - Run pilot with logistics team, T017 - Collect pilot feedback and metrics, T018 - Iterate on model based on feedback, T019 - Prepare production deployment plan, T020 - Staff training program, T021 - Full production rollout, T022 - Set up continuous monitoring, T025 - Legal review of data usage, T026 - Complete staff training rollout (remaining sites), T027 - Full production rollout sign-off (all distribution centers), T028 - Monitor model performance post-launch, T029 - Conduct final project review workshop, T030 - Publish recommendations for Phase 2 scaling Blocked: T024 - Security review of AI infrastructure (BLOCKED: waiting for security team availability) At risk: T005 - Map current logistics workflows (due date 2026-09-19 has passed, not done) Upcoming milestones: M03 - Current systems audit report, M04 - Gap analysis report, M05 - Final prioritized use case list with success criteria, M06 - Data pipelines established, M07 - Working prototype models, M08 - Pilot results report, M09 - User feedback and performance metrics, M10 - Updated models and process refinements, M11 - Full production rollout, M12 - Training sessions completed, M13 - Project closure report, M14 - Recommendations for future scaling <i>Automated with this n8n workflow</i>

Screenshots as below:

Ask Mira

Submits to your n8n workflow's webhook. Works for plan, risk, and status requests, same as the chat window.

Webhook URL

http://localhost:5678/webhook/MiraBot-newerfined

Your request

Summarize current project status: tasks done, in progress, to do, blocked?

Send request

Done

Scoped to All Sprints — 4 done, 23 to do, 1 blocked, 1 at risk, 12 milestones found.

Status report for All Sprints (overall health: Red):
Done: T001 - Define project scope and objectives, T002 - Form cross-functional task force, T003 - Initial stakeholder alignment meeting, T004 - Audit existing data infrastructure
To do: T006 - Identify data gaps and quality issues, T007 - Shortlist AI use cases, T008 - Define success criteria for each use case, T009 - Vendor evaluation for ML platform, T010 - Collect historical demand data, T011 - Data cleaning and preprocessing pipeline, T012 - Design demand forecasting model, T013 - Design route optimization model, T014 - Build model training pipeline, T015 - Deploy demand model to test environment, T016 - Run pilot with logistics team, T017 - Collect pilot feedback and metrics, T018 - Iterate on model based on feedback, T019 - Prepare production deployment plan, T020 - Staff training program, T021 - Full production rollout, T022 - Set up continuous monitoring, T023 - Legal review of data usage, T026 - Complete staff training rollout (remaining sites), T027 - Full production rollout sign-off (all distribution centers), T028 - Monitor model performance post-launch, T029 - Conduct final project review workshop, T030 - Publish recommendations for Phase 2 scaling
Blocked: T004 - Security review of AI infrastructure (BLOCKED: waiting for security team availability)
At risk: T005 - Map current logistics workflows (due date 2026-09-19 has passed, not done)
Upcoming milestones: M03 - Current systems audit report, M04 - Gap analysis report, M05 - Final prioritized use case list with success criteria, M06 - Data pipelines established, M07 - Working prototype models, M08 - Pilot results report, M09 - User feedback and performance metrics, M10 - Updated models and process refinements, M11 - Full production rollout, M12 - Training sessions completed, M13 - Project closure report, M14 - Recommendations for future scaling

Directories

Starred

Drag and drop important stuff here

Channels

errors

n8nflowaggngoldive-projectupdates

project-questions-errors

Direct messages

Agents & apps

Mira Bot

Project Updates

Connect apps

Scoped to All Sprints — 4 done, 23 to do, 1 blocked, 1 at risk, 12 milestones found.

Status report for All Sprints (overall health: Red):
Done: T001 - Define project scope and objectives, T002 - Form cross-functional task force, T003 - Initial stakeholder alignment meeting, T004 - Audit existing data infrastructure
To do: T006 - Identify data gaps and quality issues, T007 - Shortlist AI use cases, T008 - Define success criteria for each use case, T009 - Vendor evaluation for ML platform, T010 - Collect historical demand data, T011 - Data cleaning and preprocessing pipeline, T012 - Design demand forecasting model, T013 - Design route optimization model, T014 - Build model training pipeline, T015 - Deploy demand model to test environment, T016 - Run pilot with logistics team, T017 - Collect pilot feedback and metrics, T018 - Iterate on model based on feedback, T019 - Prepare production deployment plan, T020 - Staff training program, T021 - Full production rollout, T022 - Set up continuous monitoring, T023 - Legal review of data usage, T026 - Complete staff training rollout (remaining sites), T027 - Full production rollout sign-off (all distribution centers), T028 - Monitor model performance post-launch, T029 - Conduct final project review workshop, T030 - Publish recommendations for Phase 2 scaling
Blocked: T004 - Security review of AI infrastructure (BLOCKED: waiting for security team availability)
At risk: T005 - Map current logistics workflows (due date 2026-09-19 has passed, not done)
Upcoming milestones: M03 - Current systems audit report, M04 - Gap analysis report, M05 - Final prioritized use case list with success criteria, M06 - Data pipelines established, M07 - Working prototype models, M08 - Pilot results report, M09 - User feedback and performance metrics, M10 - Updated models and process refinements, M11 - Full production rollout, M12 - Training sessions completed, M13 - Project closure report, M14 - Recommendations for future scaling
Automated with this n8n workflow

Use case T11 with Chat Trigger

What milestones are coming up in the next 2 weeks?

ID	Test Input	Type	Expected Output Must Contain	Actual Output
T11	What milestones are coming up in the next 2 weeks?	Timeline	Must reference actual milestones from project_timeline.csv. Must NOT invent milestones.	Upcoming milestones in the next 2 weeks — 3 found. Status report for All Sprints (overall health: Green): Upcoming milestones: M14 - Recommendations for future scaling, M13 - Project closure report, M29 - Conduct final project review workshop <i>Automated with this n8n workflow</i>

Screenshots as below:

What milestones are coming up in the next 2 weeks?

Success in 7.667s

RTT: Flowscore Get (Status Context)

RTT: Embeddings OpenAI

OpenAI Chat Model

What milestones are coming up in the next 2 weeks?

output

Upcoming milestones in the next 2 weeks — 3 milestones found. Status report for All Sprints (overall health: Green). Upcoming milestones: M14 - Gap analysis report, M03 - Current systems audit report, M05 - Final prioritized use case list with success criteria

Project: Ignorance review

Direct messages

Agents & apps

Mira Bot

Project Updates

Connect apps

Upcoming milestones in the next 2 weeks.

Scoped to All Sprints — 3 milestones found.

Status report for All Sprints (overall health: Green):
Upcoming milestones: M14 - Gap analysis report, M03 - Current systems audit report, M05 - Final prioritized use case list with success criteria
Automated with this n8n workflow

@Mira Bot

Shift + Enter to add a new line


Use case T12 with Chat Trigger

(This Example is shown in the Observability section on Tracing and Prompt fixes done further to improve the response)

Generate a stakeholder update email for Sprint 2 progress.

ID	Test Input	Type	Expected Output Must Contain	Actual Output
T12	Generate a stakeholder update email for Sprint 2 progress.	Comms	Must reference Sprint 2 tasks. Must be professional. Must NOT include Sprint 3+ tasks.	 Sprint 2 Progress Update Hello, In Sprint 2, the audit of existing data infrastructure (task T004) has been completed successfully. The mapping of current logistics workflows (task T005) is currently in progress, while the identification of data gaps and quality issues (task T006) remains to be done. However, task T005 - Map current logistics workflows (due 2026-09-19, not yet done) has missed its deadline. Next, the team will focus on completing the analysis tasks and preparing for subsequent phases. Best regards, Mira Automated with this n8n workflow


Screenshots as below:




The image shows two side-by-side screenshots. The left screenshot displays the n8n workflow editor with a complex flowchart of nodes and connections. The right screenshot shows a Slack message from 'Mira Bot' titled 'Sprint 2 Progress Update'. The message content is identical to the 'Actual Output' column in the table above, including the email icon, greeting, progress update, and signature.

Daily Milestone Tracking – Update to Slack

Daily Milestone Tracking is automated with a trigger node run once daily and is updated to Slack channel.



Slack Result seen as below:



For Slack Approval - Human in the Loop

Use case T13 – Human in the Loop Example1

Priya Nair is leaving the organization. Which sprints does this affect?

ID	Test Input	Type	Actual Output
T12	Priya Nair is leaving the organization. Which sprints does this affect?	Risk (detailed)	⚠️ Priya Nair's departure affects Sprint 2 and Sprint 4 due to her involvement in key tasks and milestones. ⚡ Added 2 new risk(s) to the register: R01 - Loss of continuity and expertise in data management and analysis due to Priya Nair's departure. (mitigation: Reassign tasks to other team members and provide necessary training to ensure smooth transition.) R02 - Disruption in the execution of tasks related to data cleaning and preprocessing in Sprint 4. (mitigation: Identify and onboard a replacement for Priya Nair as soon as possible to maintain project momentum.) Automated with this n8n workflow

Screenshots as below:

RISK REGISTER BEFORE UPDATION: →

risk_id	category	risk_challenge	impact	mitigation_strategy
R01	Data Quality & Availability	Incomplete or unstructured data	Poor model performance and inaccurate outcomes	Perform early data audits and implement robust data governance
R02	Integration Issues	Difficulty integrating AI with legacy systems	Delays in deployment and operational disruptions	Plan integration roadmap and use API/intermediaries
R03	Talent & Skills Gap	Lack of internal AI/ML expertise	Overreliance on external vendors and slow progress	Provide staff training and hire specialized talent
R04	Change Resistance	Employees resist adopting AI-driven workflows	Low adoption and reduced ROI	Strong change management and early engagement
R05	Ethical & Legal Concerns	Risks of biased models or data privacy violations	Legal repercussions and reputational damage	Involve legal early and implement fairness checks
R06	Cost Overruns	Underestimation of AI development costs	Budget constraints and project delays	Set realistic budget and build in contingency funds
R07	Unrealistic Expectations	Stakeholders expect immediate ROI	Disillusionment and loss of support	Align expectations and show quick wins
R08	Model Performance Drift	AI models lose accuracy over time	Incorrect decisions and trust erosion	Regular retraining cycles and monitoring dashboards
R09	Security Vulnerabilities	AI models targeted by adversarial attacks	Data breaches and operational risks	Adopt AI security best practices and penetration testing

Waiting for Slack Approval for addition of new Risks

Waiting | 236KB | ID#785 | Autosave

+ Add tag

Chat

Priya Nair is leaving the organization. Which sprints does this affect?

Logs

Waiting for 1m 16s 16,856 Tokens

Success in 1.609s 1,023 Tokens

INPUT

N1: When chat message received

N2: User Input

N3: Classifier Agent

OUTPUT

chatinput Priya Nair is leaving the organization. Which sprints does this affect?

sessionId edff2b6f21fb4861af042076safaf33a

source chat

Mira Bot APP 1:04 PM

Approve or Deny to Append the following New Risks to Risk Register:

2 New risks found:

Loss of continuity in data analysis and pipeline development due to Priya Nair's departure. (Mitigation: Reassign tasks to other team members and ensure knowledge transfer before her departure.)

Potential delays in achieving milestones related to data pipelines and analysis. (Mitigation: Identify and onboard a replacement or redistribute tasks among existing team members.)

Deny Approve

Final Result

+ Add tag

Chat

Priya Nair is leaving the organization. Which sprints does this affect?

Logs

Success in 3m 19.469s 16,606 Tokens

N23: Confirm Risk Added

Final Gate

Post Message to Slack

INPUT

output Priya Nair's departure affects Sprint 2 and Sprint 4 due to her involvement in key tasks and milestones. Added 2 new risk(s) to the register: - R01 - Loss of continuity in data analysis and pipeline development due to Priya Nair's departure. (Mitigation: Reassign tasks

RISK REGISTER AFTER UPDATION: →

risk_id	category	risk_challenge	impact
R01	Data Quality & Availability	Incomplete or unstructured data	Poor model perform
R02	Integration Issues	Difficulty integrating AI with legacy systems	Delays in deployment
R03	Talent & Skills Gap	Lack of internal AI/ML expertise	Overreliance on ex
R04	Change Resistance	Employees resist adopting AI-driven workflows	Low adoption and r
R05	Ethical & Legal Concerns	Risks of biased models or data privacy violations	Legal repercussions
R06	Cost Overruns	Underestimation of AI development costs	Budget constraints
R07	Unrealistic Expectations	Stakeholders expect immediate ROI	Disillusionment anc
R08	Model Performance Drift	AI models lose accuracy over time	Incorrect decisions
R09	Security Vulnerabilities	AI models targeted by adversarial attacks	Data breaches and
R10	Resource Bottleneck	Loss of continuity in data analysis and pipeline developm	Delays in project tir
R11	Milestone Delays	Potential delays in achieving milestones related to data pi	Impact on project ti

With details into the risk register added as below for the following columns:

SLACK Result Post Approval – R10 and R11 risks newly added.

Mira Bot APP 1:04 PM

Approve or Deny to Append the following New Risks to Risk Register:

2 New risks found:

Loss of continuity in data analysis and pipeline development due to Priya Nair's departure. (Mitigation: Reassign tasks to other team members and ensure knowledge transfer before her departure.)

Potential delays in achieving milestones related to data pipelines and analysis. (Mitigation: Identify and onboard a replacement or redistribute tasks among existing team members.)

Deny Approve

Priya Nair's departure affects Sprint 2 and Sprint 4 due to her involvement in key tasks and milestones.

Added 2 new risk(s) to the register:

R10 - Loss of continuity in data analysis and pipeline development due to Priya Nair's departure. (Mitigation: Reassign tasks to other team members and ensure knowledge transfer before her departure.)

R11 - Potential delays in achieving milestones related to data pipelines and analysis. (Mitigation: Identify and onboard a replacement or redistribute tasks among existing team members.)

Automated with this n8n workflow

risk_id, category, risk_challenge, impact, mitigation_strategy, last_updated					
R11	Milestone Delays	Potential delays in achieving milestones related to data pipelines and analysis.	Impact on project timelines and deliverables for Sprint 2 and Sprint 4.	Identify and onboard a replacement or redistribute tasks among existing team members.	2026-09-21

Use case T14 – Human in the Loop Example1

Arjun Mehta might be stepping back from the project soon, what's the impact?

ID	Test Input	Type	Actual Output
T12	Arjun Mehta might be stepping back from the project soon, what's the impact?	Risk (detailed)	⚠️ Arjun Mehta's potential departure could significantly impact several upcoming sprints, particularly Sprint 6, Sprint 9, and Sprint 12, where he is assigned critical tasks. ⚡ Added 1 new risk(s) to the register: R12 - Loss of continuity in project tasks due to Arjun Mehta's departure. (mitigation: Reassign tasks to other team members and ensure knowledge transfer

Screenshots as below:


RISK REGISTER BEFORE UPDATION: →

risk_id	category	risk_challenge	impact
R01	Data Quality & Availability	Incomplete or unstructured data	Poor model performance and inaccurate output
R02	Integration Issues	Difficulty integrating AI with legacy systems	Delays in deployment and operational disruption
R03	Talent & Skills Gap	Lack of internal AI/ML expertise	Overreliance on external vendors and slow progress
R04	Change Resistance	Employees resist adopting AI-driven workflows	Low adoption and reduced ROI
R05	Ethical & Legal Concerns	Risks of biased models or data privacy violations	Legal repercussions and reputational damage
R06	Cost Overruns	Underestimation of AI development costs	Budget constraints and project delays
R07	Unrealistic Expectations	Stakeholders expect immediate ROI	Disillusionment and loss of support
R08	Model Performance Drift	AI models lose accuracy over time	Incorrect decisions and trust erosion
R09	Security Vulnerabilities	AI models targeted by adversarial attacks	Data breaches and operational risks
R10	Resource Bottleneck	Loss of continuity in data analysis and pipeline development	Delays in project timelines and potential quality issues
R11	Milestone Delays	Potential delays in achieving milestones related to data pipelines and analysis	Impact on project timelines and deliverables for

Waiting for Slack Approval for addition of new Risks

Waiting | 360KB | ID#788 | Autosave

+ Add tag



Chat

Arjun Mehta might be stepping back from the project soon, what's the impact?

Logs

Waiting for 1m 17s | 36,387 Tokens

N1: When chat message received

N2: User Input

N3: Classifier Agent

Success in 1.85s | 1,029 Tokens

INPUT

chatinput: Arjun Mehta might be stepping back from the project soon, what's the impact?

sessionId: edf22b6f21fb4861af042f076faf33a

source: chat

Output

1 item

Approve or Deny to Append the following New Risks to Risk Register:

1 New risks found:

Loss of continuity in project tasks due to Arjun Mehta's departure. (mitigation: Reassign tasks to other team members and ensure knowledge transfer before his departure.)

Deny Approve

Final Result

LangSmith - Tracing Home Page

LangSmith

Application

Mira-AgenticApp

Search

Home

Tracing

Engine

Monitoring

Datasets & Experiments

Evaluators

Annotation Queues

Custom Apps

Prompts

Playground

Studio

Context Hub

Deployments

Sandboxes

LLM Gateway

Trace retention is changing

Reminder: Beginning September 14, 2026, the maximum retention period for LangSmith traces has changed to 180 days.

Learn more

Mira-AgenticApp / Tracing

Project

Search by name...

Columns

Name	Most Recent Run (7d)	Trace Count (7d)	Error Rate (7d)	P50 Latency (7d)	P99 Latency (7d)	Total Tokens (7d)	Total Cost (7d)	Description
MiraAgenticSolution	23 minutes ago	587	3%	148s	109.40s	3.076M	\$0.6323	
evaluators	19 hours ago	282	0%	2.17s	6.52s	206.7K	\$0.1517	

LangSmith

Application

Mira-AgenticApp

Search

Home

Tracing

Engine

Monitoring

Datasets & Experiments

Evaluators

Annotation Queues

Custom Apps

Prompts

Playground

Studio

Context Hub

Deployments

Sandboxes

LLM Gateway

Trace retention is changing

Reminder: Beginning September 14, 2026, the maximum retention period for LangSmith traces has changed to 180 days.

Learn more

Mira-AgenticApp / Tracing / MiraAgenticSolution

Project

Search by name...

Columns

Name	Most Recent Run (7d)	Trace Count (7d)	Error Rate (7d)	P50 Latency (7d)	P99 Latency (7d)	Total Tokens (7d)	Total Cost (7d)	Description
MiraAgenticSolution	23 minutes ago	587	3%	148s	109.40s	3.076M	\$0.6323	
evaluators	19 hours ago	282	0%	2.17s	6.52s	206.7K	\$0.1517	

Trace

Waterfall

Filter runs...

[Mira-AIPoweredIntelligentAssistant-NewRefined] N2S: Status Reporter Agent 12:06 PM

ChatOpenAI gpt-4o-mini 3.77s 10.62K

ToolCallingAgentOutputParser 0.00s

Some runs have been hidden. Show 8 hidden runs

Trace

Feedback

Input

Output

Attributes

Feedback

Add feedback

Input

User

Today's date: 2026-09-21

User request: Generate a weekly status report using the sample task board for Sprint 3.

CRITICAL STEP FOR STATUS REPORTING:

1. Cross-examine the "User request" against the actual task boards, timelines, and sprint logs retrieved in your context window.

2. Every count, task ID, and status metric must be pulled verbatim from the files. Do not estimate, combine separate metrics, or round down numbers.

3. If tasks are flagged with dependencies or blocked indicators, treat those structural conflicts as explicit log requirements in the final report output.

Additional Fields

Output

Fields

log { "sufficient_data": true, "no_actionable_request": false, "report_scope": "Sprint 3", "answer": "", "requested_category": ...

returnValues

output

{ "sufficient_data": true, "no_actionable_request": false, "report_scope": "Sprint 3", "answer": "", "requested_categories": ["all"], "tasks": {

{

12:06 PM 12:06 PM

LangSmith - Datasets and Experiments

LangSmith

Application

Mira-AgenticApp

Search

Home

Tracing

Engine

Monitoring

Datasets & Experiments

Evaluators

Annotation Queues

Custom Apps

Prompts

Playground

Studio

Context Hub

Deployments

Sandboxes

LLM Gateway

Trace retention is changing

Reminder: Beginning September 14, 2026, the maximum retention period for LangSmith traces has changed to 180 days.

Learn more

Mira-AgenticApp / Datasets & Experiments

+ Dataset + Experiment

Search by name...

Columns

	Dataset	IT	Description	Experiments	IT	Last Experiment	Examples	Updated At	IT	Created At	IT
	☐ mira_Ds_v5		Evaluation Dataset with Columns id, agent, ty...	1		19 hours ago	12	19 hours ago		19 hours ago	
	☐ mira_Ds_v4		Evaluation Dataset with the Column details as...	1		20 hours ago	12	20 hours ago		20 hours ago	
	☐ mira_Ds_v3		Dataset Details with columns: id, agent (agent...	3		20 hours ago	10	4 days ago		4 days ago	
	☐ mira_DS-ipoptsdosdnts		Dataset Details with columns: id, agent (agent ...	1		5 days ago	10	5 days ago		5 days ago	

Evaluators run from Command Prompt against the new Dataset created

```
PS C:\users\srvividya\appdata\Mira-finetuning> python CheckforDatasetID.py
Datasets visible to this API key:
- mira_DS_v4 (id: 82ff7ad6c-bbcb-4bf6-b44d-a929e1ecc18b)
- mira_DS_v3 (id: 17383dad-3b35-4d97-a536-487f2ff6bfe9)
- mira_DS-ipoptsdosdnts (id: afbbf69b-6d19-43d7-b209-d06484578be6)
- Evaluator: mira_DS_v3:Conciseness (81478e1c) - conciseness (id: f50ece24-1965-4b0b-9b71-9e848481074c)
- Evaluator: mira_DS_v3:Answer Relevance (edd60ea3) - answer_relevance (id: 6755c228-49a3-4184-9adc-30100ca0b574)
- Evaluator: mira_DS_v3:Assertions (e0dde8b8) - assertions_passed (id: 3d85f67c-61a4-49a9-af8c-2b101d62dfff)
- Evaluator: mira_DS_v3:Bias & Fairness (e4c7178f) - fairness (id: 4f7846f2-43be-4889-9319-377e45aad0bc)
- Evaluator: mira_DS_v3:Hallucination (5b78306a) - hallucination (id: dea29f00-a63b-4fdd-8491-80111548888b)
- mira_DS_v5 (id: 0ae436fa-a175-4056-b2f8-8dcfdbae88d4)
PS C:\users\srvividya\appdata\Mira-finetuning> python LangSmith-Eval-Python_script.py
Detected dataset input key: 'id'
Detected dataset output key: 'input'

View the evaluation results for experiment: 'mira-eval-allTsusecases-ae34171e' at:
https://smith.langchain.com/o/75f0824e-37d5-41b9-a793-292f5cbd95b9/datasets/0ae436fa-a175-4056-b2f8-8dcfdbae88d4/compare?selectedSessions=7f093859-2d21-4daf-bbe5-39d2afc85c4c

12it [00:00, 25.77it/s]

Evaluation Results:
<ExperimentResults mira-eval-allTsusecases-ae34171e>
```

LangSmith

Application

Mira-AgenticApp

Search

Home

Tracing

Engine

Monitoring

Datasets & Experiments

Evaluators

Annotation Queues

Custom Apps

Prompts

Playground

Studio

Context Hub

Deployments

Sandboxes

LLM Gateway

Trace retention is changing

Reminder: Beginning September 14, 2026, the maximum retention period for LangSmith traces has changed to 180 days.

Learn more

Mira-AgenticApp / Datasets & Experiments / mira_Ds_v5

All Versions + Evaluator + Experiment

Search by name...

Group by Filters

Feedback

Latency

Tokens

Cost

Experiment	IT	Splits	Repetitions	Progress	answer_f...	IT	concisen...	IT	coverage	IT	fairness	IT	hallucinat...	IT	Created At
#1	mira-eval-allTsusecases-ae3417...	base	1	12 / 12 runs	0.75		0.74		1.00		0.00		0.08		9/20/2026, 5

Page 1 Show 10

LangSmith

Trace retention is changing
Reminder: Beginning September 14, 2026, the maximum retention period for LangSmith traces has changed to 180 days.

Mira-AgenticApp / Datasets & Experiments / mira_Ds_v5 / mira-eval-allTausecases-ae3471e

mira-eval-allTausecases-ae3471e 12 / 12 runs

Inputs	Reference Outputs	Outputs	conc...	coverage	fairm...	hallucination	Latency
1 N11: Planner Agent	I need a bit more detail to generate a plan.	Generate a project plan for a 2-week proje	1.00	1.00	0.00	1.00	0.14s
2 N25: Status Reporter Agent	Status report for All Sprints (overall health:	Summarize current project status: tasks d	1.00	1.00	0.00	0.00	0.00s
3 N30: Stakeholder Update Agent	Sprint 2 Progress Update: Current State An	Generate a stakeholder update email for S	1.00	1.00	0.00	0.00	0.00s
4 N25: Status Reporter Agent	Blocked tasks include T024 - Security revi	Which tasks are blocked or at risk of miss	1.00	1.00	0.00	0.00	0.00s
5 N16: Risk Assessor Agent	Here are the top 3 risks for the ABCDE Ltd	What are the top 3 risks for the ABCDE Ltd	1.00	1.00	0.00	0.00	0.00s
6 N11: Planner Agent	I need a bit more detail to generate a plan.	Generate a project plan for: 'We want to b	0.00	1.00	0.00	0.00	0.00s
7 N25: Status Reporter Agent	Could you share more details on what you'	Generate a weekly status report for: 'Thin	0.60	1.00	0.00	0.00	0.00s
8 N25: Status Reporter Agent	Upcoming milestones in the next 2 weeks -	What milestones are coming up in the nex	1.00	1.00	0.00	0.00	0.00s
9 N25: Status Reporter Agent	Scoped to Sprint 3 — 4 tasks found. Status	Generate a weekly status report using the	1.00	1.00	0.00	0.00	0.00s
10 N11: Planner Agent	I am generating a project plan for the AI Ad	Generate a project plan for the AI Adoptio	1.00	1.00	0.00	0.00	0.00s
11 N16: Risk Assessor Agent	Here are the relevant risks for the AI Adopt	Generate a risk assessment for the AI Ado	0.30	1.00	0.00	0.00	0.00s
12 N16: Risk Assessor Agent	I need more detail to assess risk properly. I-	Generate a risk assessment for: 'New proj	0.00	1.00	0.00	0.00	0.00s

Langsmith - Datasets and Experiments – Offline Tracing

Example Fix1 – [Use Case T10]

Use case: Summarize current project status: tasks done, in progress, to do, blocked?

Issue: Erroneous output and giving undefined-undefined placeholders in the output. Updated prompt to fix further.

LangSmith

Trace retention is changing
Reminder: Beginning September 14, 2026, the maximum retention period for LangSmith traces has changed to 180 days.

Mira-AgenticApp / Datasets & Experiments / mira_Ds_v5 / mira-eval-backfill-a168fa9e

mira-eval-backfill-a168fa9e 10 / 10 runs

Target 9:36 AM 0.00s

Inputs	Reference Outputs
1 What are the top 3 risks for the ABCDE L...	Here are the top 3 risks for
2 Generate a risk assessment for: 'New pr...	I'm not sure what you're lo
3 Summarize current project status: tasks ...	Status report for All Sprints
4 Generate a weekly status report for: 'Thin...	Could you share more deta
5 Generate a project plan for: 'We want to ...	I need a bit more detail to g
6 Which tasks are blocked or at risk of mis...	Scoped to all tasks — 1 blo
7 Generate a project plan for a 2-week pro...	I need a bit more detail to g
8 Generate a project plan for the AI Adopt...	I am generating a project p
9 Generate a risk assessment for the AI Ad...	Here are the relevant risks
10 Generate a weekly status report using th...	Scoped to Sprint 3 — 4 tas

Feedback Input Output Reference Outputs Attributes

Reference Outputs

Fields

- actual_reference_output: Status report for All Sprints (overall health: Amber): Done: T001 - Define project scope and objectives, TO...
- conclusions_expectation: Comprehensive but organized by status category; should not include placeholder/undefined entries.
- context: «full_task_board» - T001 | Define project scope and objectives | Status: Done | Assignee: Arjun Mehta | Sprint: Sprint 1 | D...
- entity_pool: T001;T002;T003;T004;T005;T006;T007;T008;T009;T010;T011;T012;T013;T014;T015;T016;T017;T018;T019;T020;T021;T...
- expected_deliverable: Full status breakdown across Done / In Progress / To Do / Blocked, grounded in the task board
- known_reference_output_flags: This captured output contains 5 literal 'undefined - undefined' placeholder entries and misclassifies T005 as 'In progress' though the Task Board marks it 'Done'. Known data-quality issue --- useful for verifying Hallucination/Groundedness evaluators correctly flag it.
- must_contain: Must count correctly from task board CSV. Counts must match actual data.
- must_not_contain: Must NOT round or estimate counts.
- scenarios_type: grounded
- key_facts_expected: null

Attributes

Metadata

- __ls_runner: pyc.sdk.evaluate
- example_version: 2026-09-07T01:59:42:12:4444-00:00
- num_repetitions: 1
- revision_id
- LangSmith:
 - LANGSMITH_ENDPOINT: https://api.smith.langchain.com/
 - LANGSMITH_PROJECT: MiraAgenticSolution
 - LANGSMITH_TRACING: false
 - ls_example_agent: N25: Status Reporter Agent
 - ls_example_type: Status summary

Use case Status	ID	Test Input	Type	Expected Output Must Contain	Actual Output
Before	T10	Summarize	Status	Must count	Status report for All Sprints (overall health: Amber):

Tracing/ Fixing the Prompt		current project status: tasks done, in progress, to do, blocked?	summary (Observed Hallucination in Langsmith)	correctly from task board CSV. Counts must match actual data.	Done: T001 - Define project scope and objectives, T002 - Form cross-functional task force, T003 - Initial stakeholder alignment meeting, T004 - Audit existing data infrastructure, undefined - undefined, undefined - undefined, undefined - undefined To do: T006 - Identify data gaps and quality issues, T007 - Shortlist AI use cases, T008 - Define success criteria for each use case, T009 - Vendor evaluation for ML platform, T010 - Collect historical demand data, T011 - Data cleaning and preprocessing pipeline, T012 - Design demand forecasting model, T013 - Design route optimization model, T014 - Build model training pipeline, T015 - Deploy demand model to test environment, T016 - Run pilot with logistics team, T017 - Collect pilot feedback and metrics, T018 - Iterate on model based on feedback, T019 - Prepare production deployment plan, T020 - Staff training program, T021 - Full production rollout, T022 - Set up continuous monitoring, T023 - Project closure report, T024 - Security review of AI infrastructure, T025 - Legal review of data usage, T026 - Complete staff training rollout (remaining sites), T027 - Full production rollout sign-off (all distribution centers), T028 - Monitor model performance post-launch, T029 - Conduct final project review workshop, T030 - Publish recommendations for Phase 2 scaling Blocked: T024 - Security review of AI infrastructure (BLOCKED: waiting for security team availability) At risk: T005 - Map current logistics workflows, undefined - undefined, undefined - undefined Automated with this n8n workflow
After Tracing/ Fixing the Prompt	T10	Summarize current project status: tasks done, in progress, to do, blocked?	Status summary (Observed Hallucination in Langsmith)	Must count correctly from task board CSV. Counts must match actual data.	 Scoped to All Sprints — 4 done, 23 to do, 1 blocked, 1 at risk, 12 milestones found.  Status report for All Sprints (overall health: Red): Done: T001 - Define project scope and objectives, T002 - Form cross-functional task force, T003 - Initial stakeholder alignment meeting, T004 - Audit existing data infrastructure To do: T006 - Identify data gaps and quality issues, T007 - Shortlist AI use cases, T008 - Define success criteria for each use case, T009 - Vendor evaluation for ML platform, T010 - Collect historical demand data, T011 - Data cleaning and preprocessing pipeline, T012 - Design demand forecasting model, T013 - Design route optimization model, T014 - Build model training pipeline, T015 - Deploy demand model to test environment, T016 - Run pilot with logistics team, T017 - Collect pilot feedback and metrics, T018 - Iterate on model based on feedback, T019 - Prepare production deployment plan, T020 - Staff training program, T021 - Full production rollout, T022 - Set up continuous monitoring, T025 - Legal review of data usage, T026 - Complete staff training rollout (remaining sites), T027 - Full production rollout sign-off (all distribution centers), T028 - Monitor model performance post-launch, T029 - Conduct final project review workshop, T030 - Publish recommendations for Phase 2 scaling Blocked: T024 - Security review of AI infrastructure (BLOCKED: waiting for security team availability) At risk: T005 - Map current logistics workflows (due date 2026-09-19 has passed, not done) Upcoming milestones: M03 - Current systems audit report, M04 - Gap analysis report, M05 - Final prioritized use case list with success criteria, M06 - Data pipelines established, M07 - Working prototype models, M08 - Pilot results report, M09 - User feedback and performance metrics, M10 - Updated models and process refinements, M11 - Full production rollout, M12 - Training sessions completed, M13 - Project closure report, M14 - Recommendations for future scaling Automated with this n8n workflow

ExampleFix2 – [Use Case T12]

Use case: Generate a stakeholder update email for Sprint 2 progress.

Issue: Erroneous output for wrongly stating the Task status. Updated prompt to fix further.

Use case Status	Test Input	Type	Expected Output Must Contain	Actual Output
Before Tracing/ Fixing the Prompt	Generate a stakeholder update email for Sprint 2 progress.	Comms	Must reference Sprint 2 tasks. Must be professional. Must NOT include Sprint 3+ tasks. NOTE: It shows NO blockers. however, T005 is passed its deadline for delivery which should have been highlighted here.. This was observed further node by node at Langsmith Traces. (Observed Answer_relevance in Langsmith)	 Sprint 2 Progress Update Hello, In Sprint 2, we completed the audit of the existing data infrastructure (Task T004) and are currently progressing on mapping the current logistics workflows (Task T005). The task to identify data gaps and quality issues (Task T006) is still to be started. There are no reported blockers at this time. Next, we will focus on completing the gap analysis report and preparing for subsequent sprint activities. Best regards, Mira Automated with this n8n workflow

After Tracing/ Fixing the Prompt	Generate a stakeholder update email for Sprint 2 progress.	Comms	Must reference Sprint 2 tasks. Must be professional. Must NOT include Sprint 3+ tasks. NOTE: Issue is fixed now after prompt updates. (Observed Answer_relevance in Langsmith)	 Sprint 2 Progress Update Hello, In Sprint 2, the audit of existing data infrastructure (task T004) has been completed successfully. The mapping of current logistics workflows (task T005) is currently in progress, while the identification of data gaps and quality issues (task T006) remains to be done. However, task T005 - Map current logistics workflows (due 2026-09-19, not yet done) has missed its deadline. Next, the team will focus on completing the analysis tasks and preparing for subsequent phases. Best regards, Mira <i>Automated with this n8n workflow</i>
-------------------------------------	--	-------	---	--

Fine Tuning

Running the Scripts

Python file update with Pipeline Outputs of Risk Agent Node (N16: Risk Assessor Agent) to Risk Assessor file before running as below.

@windows Powershell

```
python test_risk_assessor.py
```

```
Windows PowerShell
PS C:\Users\srividya\AppData\Local\Microsoft\Windows\Terminal\Profiles\Windows PowerShell> python test_risk_assessor.py
C:\Users\srividya\AppData\Local\Microsoft\Windows\Terminal\Profiles\Windows PowerShell> DeprecationWarning: 'LLMTestCaseParams' is deprecated and will be removed in a future release. Use 'SingleTurnParams' instead.
from deepeval.test_case import LLMTestCase, LLMTestCaseParams

=====
MIRA RISK ASSESSOR AGENT (N16) - DeepEval Evaluation
=====

Validating pasted N16 outputs...
T1_vague_statement: OK
T2_vague_statement_alt: OK
T3_broad_named_project: OK
T4_top_n: OK
T5_top_n_alt: OK
T6_named_person_new: OK
T7_named_person_existing: OK
T8_bare_vague: OK
T9_sprint_scoped: OK

=====
EVALUATING VAGUE STATEMENTS (must flag as insufficient)
=====
★ You're running DeepEval's latest Vague Statement Handling [GVal] Metric! (using gpt-4.1, strict=False, async_mode=True)...

=====
DeepEval Evaluation Results
=====
test_case_0 (Passed 1 metrics)
test_case_1 (Passed 1 metrics)
test_case_2 (Failed 1 metrics)

Input: risks?
Actual Output:
{
  "intent": "question",
  "answer": "Here are the relevant risks for your query.",
  "requested_count": null,
  "risk_list_heading": "Here are the relevant risks for your query.",
  "sufficient_detail": true,
  "missing_info": [],
  "risks": [
    {
      "risk_id": "R001",
      "category": "Data Quality & Availability",
      "risk": "Incomplete or unstructured data",
      "impact": "Poor model performance and inaccurate outcomes",
      "mitigation": "Perform early data audits and implement robust data governance",
      "matched_entity": "project-wide assessment"
    },
    {
      "risk_id": "R002",
      "category": "Integration Issues",
      "risk": "Difficulty integrating AI with legacy systems",
      "impact": "Delays in deployment and operational disruptions",
      "mitigation": "Plan integration roadmap and use APIs/middleware",
      "matched_entity": "project-wide assessment"
    },
    {
      "risk_id": "R003",
      "category": "Talent & Skills Gap",
      "risk": "Lack of internal AI/ML expertise",
      "impact": "Overreliance on external vendors and slow progress",
      "mitigation": "Provide staff training and hire specialized talent",
      "matched_entity": "project-wide assessment"
    },
    {
      "risk_id": "R004",
      "category": "Change Resistance",
      "risk": "Employees resist adopting AI-driven workflows",
      "impact": "Slow adoption and reduced ROI",
      "mitigation": "Strong change management and early engagement",
      "matched_entity": "project-wide assessment"
    }
  ]
}
```

```
{
  "risk_id": "R005",
  "category": "Ethical & Legal Concerns",
  "risk": "Risks of biased models or data privacy violations",
  "impact": "Legal repercussions and reputational damage",
  "mitigation": "Involve legal early and implement fairness checks",
  "matched_entity": "project-wide assessment"
},
{
  "risk_id": "R006",
  "category": "Cost Overruns",
  "risk": "Underestimation of AI development costs",
  "impact": "Budget constraints and project delays",
  "mitigation": "Set realistic budget and build in contingency funds",
  "matched_entity": "project-wide assessment"
},
{
  "risk_id": "R007",
  "category": "Unrealistic Expectations",
  "risk": "Stakeholders expect immediate ROI",
  "impact": "Disillusionment and loss of support",
  "mitigation": "Align expectations and show quick wins",
  "matched_entity": "project-wide assessment"
},
{
  "risk_id": "R008",
  "category": "Model Performance Drift",
  "risk": "AI models lose accuracy over time",
  "impact": "Incorrect decisions and trust erosion",
  "mitigation": "Regular retraining cycles and monitoring dashboards",
  "matched_entity": "project-wide assessment"
},
{
  "risk_id": "R009",
  "category": "Security Vulnerabilities",
  "risk": "AI models targeted by adversarial attacks",
  "impact": "Data breaches and operational risks",
  "mitigation": "Adopt AI security best practices and penetration testing",
  "matched_entity": "project-wide assessment"
},
{
  "risk_id": "R010",
  "category": "Vendor Lock-in",
  "risk": "Overdependence on a specific AI platform",
  "impact": "Flexibility loss and future switching costs",
  "mitigation": "Favor open standards and negotiate flexible vendor terms",
  "matched_entity": "project-wide assessment"
},
{
  "risk_id": "R011",
  "category": "Talent & Skills Gap",
  "risk": "Departure of Arjun Mehta could lead to a critical gap in leadership for ongoing tasks and milestones.",
  "impact": "Potential delays in task execution and project timelines, especially for tasks involving high-level stakeholder engagement and strategic project closure.",
  "mitigation": "Identify and groom a successor from the team, redistribute tasks among existing team members to ensure continuity, and document critical processes and decisions.",
  "matched_entity": "Arjun Mehta"
}
],
  "new_risks": []
}
```

Metrics Status	Metric	Score	Threshold	Reason
FAIL	Vague Statement Handling [GEval]	0.12	0.50	The input is a vague statement with no identifiable project name or scope, but the output incorrectly sets 'sufficient_detail' to true, populates the 'risks' array with project-specific risks, and does not ask the user for more information. This fails to align with the evaluation steps, except for not populating 'new_risks'.

Aggregate Metrics

Metric	Average Score	Pass Rate	Total
Vague Statement Handling [GEval]	0.04	66.67% passed=2 failed=1	3

⚠️WARNING: No hyperparameters logged.

» Log hyperparameters to attribute prompts and models to your test runs.

✓ Evaluation completed #1 (time taken: 7.35s | token cost: 0.01062 USD)

» Test Results (3 total tests):

» Pass Rate: 66.67% | Passed: 2 | Failed: 1

» Want to share evals with your team, or a place for your test cases to live? ❤️ 🚀

» Run `deepeval view` to analyze and save testing results on Confidential AI.

EVALUATING BROAD / NAMED-PROJECT REQUESTS

🌟 You're running DeepEval's latest **Risk Groundedness [GEval]** Metric (using gpt-4.1, strict=False, async_mode=True)...

🔥 DeepEval Evaluation Results

✅ test_case_0 (Passed 1 metrics)

✅ test_case_1 (Passed 1 metrics)

Aggregate Metrics

Metric	Average Score	Pass Rate	Total
Risk Groundedness [GEval]	0.89	100.00% passed=2 failed=0	2

⚠️WARNING: No hyperparameters logged.

» Log hyperparameters to attribute prompts and models to your test runs.

✓ Evaluation completed #1 (time taken: 5.74s | token cost: 0.009208000000000001 USD)

» Test Results (2 total tests):

» Pass Rate: 100.0% | Passed: 2 | Failed: 0

» Want to share evals with your team, or a place for your test cases to live? ❤️ 🚀

» Run `deepeval view` to analyze and save testing results on Confidential AI.

EVALUATING COUNT-SPECIFIC REQUESTS (top-N)

🌟 You're running DeepEval's latest **Requested Count Enforcement [GEval]** Metric (using gpt-4.1, strict=False, async_mode=True)...

🌟 You're running DeepEval's latest **Risk Groundedness [GEval]** Metric (using gpt-4.1, strict=False, async_mode=True)...

🔥 DeepEval Evaluation Results

test_case_0 (Passed 2 metrics)

test_case_1 (Passed 2 metrics)

Aggregate Metrics

Metric	Average Score	Pass Rate	Total
Requested Count Enforcement [GEval]	1.00	100.00% passed=2 failed=0	2
Risk Groundedness [GEval]	0.87	100.00% passed=2 failed=0	2

WARNING: No hyperparameters logged.

Log hyperparameters to attribute prompts and models to your test runs.

Evaluation completed #1 (time taken: 6.81s | token cost: 0.013512 USD)

Test Results (2 total tests):

Pass Rate: 100.0% | Passed: 2 | Failed: 0

Want to share evals with your team, or a place for your test cases to live?

Run `deepeval view` to analyze and save testing results on Confidential AI.

EVALUATING NEW RISK (named person, not yet tracked)

You're running DeepEval's latest New vs Existing Risk Routing [GEval] Metric (using gpt-4.1, strict=False, async_mode=True)...

You're running DeepEval's latest Risk Groundedness [GEval] Metric (using gpt-4.1, strict=False, async_mode=True)...

DeepEval Evaluation Results

test_case_0

Input:

Priya Nair is leaving the organization. Which sprints does this affect?

Actual Output:

```
{
  "intent": "both",
  "answer": "Priya Nair's departure will affect Sprint 2 and Sprint 4 due to her involvement in critical tasks and milestones.",
  "requested_count": null,
  "risk_list_heading": "Here are the risks affecting the sprints due to Priya Nair's departure.",
  "sufficient_detail": true,
  "missing_info": [],
  "risks": [],
  "new_risks": [
    {
      "risk_id": "",
      "category": "Talent & Skills Gap",
      "risk": "Priya Nair's departure could lead to a critical gap in leadership for ongoing tasks and milestones.",
      "impact": "Potential delays in task execution and project timelines, especially for tasks involving data analysis and pipeline development.",
      "mitigation": "Identify and groom a successor from the team, redistribute tasks among existing team members to ensure continuity, and document critical processes and decisions."
    }
  ]
}
```

Metrics

Status	Metric	Score	Threshold	Reason
PASS	New vs Existing Risk Routing [GEval]	1.00	0.50	The input describes a real, specific event involving a named individual and affected sprints. However, the 'risks' array is empty, and the risk is instead placed in 'new_risks', which is not standard per the evaluation steps. The risk entry itself is plausible and specific, but the lack of proper placement and the absence of risk_id reduces alignment. There are no unrelated or generic risks, but the structure and completeness are lacking.
FAIL	Risk Groundedness [GEval]	0.41	0.50	

Aggregate Metrics

Metric	Average Score	Pass Rate	Total
New vs Existing Risk Routing [GEval]	1.00	100.00% passed=1 failed=0	1
Risk Groundedness [GEval]	0.41	0.00% passed=0 failed=1	1

WARNING: No hyperparameters logged.

Log hyperparameters to attribute prompts and models to your test runs.

Evaluation completed #1 (time taken: 7.29s | token cost: 0.006919999999999999 USD)

Test Results (1 total tests):

Pass Rate: 0.0% | Passed: 0 | Failed: 1

Want to share evals with your team, or a place for your test cases to live?

Run `deepeval view` to analyze and save testing results on Confidential AI.

EVALUATING EXISTING RISK MATCH (named person, already tracked)

You're running DeepEval's latest New vs Existing Risk Routing [GEval] Metric (using gpt-4.1, strict=False, async_mode=True)...

You're running DeepEval's latest Risk Groundedness [GEval] Metric (using gpt-4.1, strict=False, async_mode=True)...

DeepEval Evaluation Results

test_case_0 (Passed 2 metrics)

Aggregate Metrics

Metric	Average Score	Pass Rate	Total
New vs Existing Risk Routing [GEval]	0.99	100.00% passed=1 failed=0	1
Risk Groundedness [GEval]	1.00	100.00% passed=1 failed=0	1

WARNING: No hyperparameters logged.

Log hyperparameters to attribute prompts and models to your test runs.

Evaluation completed #1 (time taken: 5.64s | token cost: 0.006954 USD)

Test Results (1 total tests):

Pass Rate: 100.0% | Passed: 1 | Failed: 0

Want to share evals with your team, or a place for your test cases to live? Run `deepeval view` to analyze and save testing results on Confidential AI.

EVALUATION COMPLETE

What to look for:

Vague statements: Vague Statement Handling should PASS

Broad/named: Risk Groundedness should PASS

Top/N: Requested Count Enforcement should PASS

Named person: New vs Existing Risk Routing should PASS

Save a screenshot for before/after comparison.

PS C:\Users\srividya\AppData\Local\Microsoft\Windows\PythonTestRunner> python test_risk_assessor.py

C:\Users\srividya\AppData\Local\Microsoft\Windows\PythonTestRunner> python test_risk_assessor.py:43: DeprecationWarning: 'LLMTestCaseParams' is deprecated and will be removed in a future release. Use 'SingleTurnParams' instead.

from deepeval.test_case import LLMTestCase, LLMTestCaseParams

MIRA RISK ASSESSOR AGENT (MIRA) - DeepEval Evaluation

Validating pasted MIRA outputs...

T1_vague_statement: OK

T2_vague_statement_alt: OK

T3_broad_named_project: OK

T4_top_n: OK

T5_top_n_alt: OK

T6_named_person_new: OK

T7_named_person_existing: OK

T8_hare_vague: OK

T9_sprint_scoped: OK

EVALUATING VAGUE STATEMENTS (must flag as insufficient)

You're running DeepEval's latest Vague Statement Handling [GEval] Metric (using gpt-4.1, strict=False, async_mode=True)...

DeepEval Evaluation Results

test_case_0 (Passed 1 metrics)

test_case_1 (Passed 1 metrics)

test_case_2

Input

Actual Output:

risks?

{
 "intent": "question",
 "answer": "Here are the relevant risks for your query.",
 "requested_count": null,
 "risk_list_heading": "Here are the relevant risks for your query.",
 "sufficient_detail": true,
 "missing_info": [],
 "risks": [
 {
 "risk_id": "R001",
 "category": "Data Quality & Availability",
 "risk": "Incomplete or unstructured data",
 "impact": "Poor model performance and inaccurate outcomes",
 "mitigation": "Perform early data audits and implement robust data governance",
 "matched_entity": "project-wide assessment",
 },
 {
 "risk_id": "R002",
 "category": "Integration Issues",
 "risk": "Difficulty integrating AI with legacy systems",
 "impact": "Delays in deployment and operational disruptions",
 "mitigation": "Plan integration roadmap and use APIs/middleware",
 "matched_entity": "project-wide assessment",
 },
 {
 "risk_id": "R003",
 "category": "Talent & Skills Gap",
 "risk": "Lack of internal AI/ML expertise",
 "impact": "Overreliance on external vendors and slow progress",
 "mitigation": "Provide staff training and hire specialized talent",
 "matched_entity": "project-wide assessment",
 },
 {
 "risk_id": "R004",
 "category": "Change Resistance",
 "risk": "Employees resist adopting AI-driven workflows",
 "impact": "Low adoption and reduced ROI",
 "mitigation": "Strong change management and early engagement",
 "matched_entity": "project-wide assessment",
 },
 {
 "risk_id": "R005",
 "category": "Ethical & Legal Concerns",
 "risk": "Risks of biased models or data privacy violations",
 "impact": "Legal repercussions and reputational damage",
 "mitigation": "Involve legal early and implement fairness checks",
 "matched_entity": "project-wide assessment",
 },
 {
 "risk_id": "R006",
 "category": "Cost Overruns",
 "risk": "Underestimation of AI development costs",
 "impact": "Budget constraints and project delays",
 "mitigation": "Set realistic budget and build in contingency funds",
 "matched_entity": "project-wide assessment",
 },
 {
 "risk_id": "R007",
 "category": "Unrealistic Expectations",
 "risk": "Stakeholders expect immediate ROI",
 "impact": "Disillusionment and loss of support",
 "mitigation": "Align expectations and show quick wins",
 "matched_entity": "project-wide assessment",
 },
 {
 "risk_id": "R008",
 "category": "Model Performance Drift",
 "risk": "AI models lose accuracy over time",
 "impact": "Incorrect decisions and trust erosion",
 "mitigation": "Regular retraining cycles and monitoring dashboards",
 "matched_entity": "project-wide assessment",
 },
],
}

⚠ **WARNING:** No hyperparameters logged.

- » Log hyperparameters to attribute prompts and models to your test runs.

✓ Evaluation completed 🕒 (time taken: 0.15s | token cost: 0.010342 USD)

» Test Results (3 total tests):

- » Pass Rate: 66.67% | Passed: 2 | Failed: 1

» Want to share evals with your team, or a place for your test cases to live? ❤

- » Run "deepseval view" to analyze and save testing results on **Confident AI**.

```

WARNING: No hyperparameters logged.
> Log hyperparameters to attribute prompts and models to your test runs.

✓ Evaluation completed 🏁 (time taken: 7.57s | token cost: 0.009172 USD)
> Test Results (2 total tests):
  > Pass Rate: 100.0% | Passed: 2 | Failed: 0

> Want to share evals with your team, or a place for your test cases to live?
  > Run "deepeval viz" to analyze and save testing results on Confidential AI.

```

Aggregate Metrics				
Metric	Average Score	Pass Rate		Total
Requested Count Enforcement [GEval]	1.00	100.00%	passed=2 failed=0	2
Risk Groundedness [GEval]	0.88	100.00%	passed=2 failed=0	2

```

WARNING: No hyperparameters logged.
» Log hyperparameters to attribute prompts and models to your test runs.

=====

✓ Evaluation completed 🕒 (time taken: 5.75s | token cost: 0.012939999999999998
Test Results (2 total tests):
  » Pass Rate: 100.0% | Passed: 2 | Failed: 0

=====


» Want to share evals with your team, or a place for your test cases to live? ❤️
  » Run "deepseval view" to analyze and save testing results on Confidential AI.

```

```

❖ You're running DeepEval's latest New vs Existing Risk Routing [GEval] Metric! (using gpt-4.1, strict=False, async_mode=True)...
❖ You're running DeepEval's latest Risk Groundedness [GEval] Metric! (using gpt-4.1, strict=False, async_mode=True)...

```



Observations against the Test Cases

Test Case	User Query - Input	Metric	Score	Status	What it means
T1 (vague statement)	Generate a risk assessment for: 'New project starting soon.	Vague Statement Handling	—	✅ PASS	Correctly asked for clarification
T2 (vague, alt phrasing)	Risk check for something new we're starting	Vague Statement Handling	—	✅ PASS	Correctly asked for clarification
T3 (broad, named project)	Generate a risk assessment for the AI Adoption Project at ABCDE Ltd.	Risk Groundedness	—	✅ PASS	Risks returned are grounded and specific
T4 (top 3)	What are the top 3 risks for the ABCDE Ltd project?	Requested Count Enforcement	1.00	✅ PASS	Exactly 3 risks returned
T4 (top 3)	What are the top 3 risks for the ABCDE Ltd project?	Risk Groundedness	—	✅ PASS	Risks are real, specific entries
T5 (top 2)	Give me the top 2 risks for this project	Requested Count Enforcement	1.00	✅ PASS	Exactly 2 risks returned
T5 (top 2)	Give me the top 2 risks for this project	Risk Groundedness	—	✅ PASS	Risks are real, specific entries
T6 (Priya Nair, new risk)	Priya Nair is leaving the organization. Which sprints does this affect?	New vs Existing Risk Routing	1.00	✅ PASS	Correctly routed to new_risks with empty risk_id
T6 (same case)	Priya Nair is leaving the organization. Which sprints does this affect?	Risk Groundedness	0.36	❌ FAIL	Metric wording penalizes the empty risks array / missing risk_id — see note below
T7 (Arjun Mehta, existing risk)	Arjun Mehta might be stepping back from the project soon, what's the impact?	New vs Existing Risk Routing	1.00	✅ PASS	Correctly matched to existing tracked risk
T7 (same case)	Arjun Mehta might be stepping back from the project soon, what's the impact?	Risk Groundedness	1.00	✅ PASS	Matched risk is specific and grounded
T8 ("risks?")	risks?	Vague Statement Handling	0.11	❌ FAIL	Set sufficient_detail: true and dumped all 11 register risks instead of asking for scope
T9 (Sprint 5, scoped)	What risks affect Sprint 5?	Risk Groundedness	—	✅ PASS	Risks returned are grounded and specific

Note: Do the corrections for handling T8 use case as it has generated an erroneous output by sharing all register risks instead of seeking the scope – Prompt needs to be corrected for this fix. Fine tuning of model will not fix this problem.